

LLD Practice Platform

Candidate Helping Guide (Optional)

Purpose

This guide gives direction without prescribing the answer. Use it to think through the assignment; your final product and design decisions should still be your own.

1. Start With the Learner Problem

Do not begin by choosing React, Node.js, MongoDB, an LLM, or a diagram library. First understand why LLD practice is hard.

- How does a learner currently practice an LLD problem?
- How does the learner decide whether the solution is good?
- What happens when two valid designs look very different?
- What feedback would help the learner improve on the next attempt?
- What evidence should the platform retain from an attempt?

2. Do a Small Amount of Real Research

Spend roughly 2–3 hours. Look at a few LLD practice/interview tools, GitHub projects, articles, or community discussions. You do not need dozens of sources.

Look for	Questions
Practice workflow	How does the learner start, work, submit, and retry?
Submission	Is the solution text, code, UML, or something else?
Feedback	Is there a score, comments, rubric, reference solution, or AI feedback?
Learning loop	Does the product help the learner identify recurring weaknesses?
Gaps	What would you change or simplify?

3. Define a Narrow MVP

A strong 2-day MVP can be very small. For example: 3–5 problems, one practice flow, one submission model, one evaluation flow, feedback, and attempt history.

Do not build features simply because they sound impressive. Ask whether each feature improves the practice loop.

4. Think About the Submission

LLD can be expressed in several ways. Pick the smallest format that gives enough evidence of design quality.

Option	What it proves
Text design	Requirements, assumptions, classes, responsibilities, reasoning.
Code	Concrete implementation, interfaces, coupling, testability, behaviour.
Diagram	Relationships, structure, responsibilities, abstraction.
Combined	More evidence, but more implementation effort.

A good decision here is not “support everything”; it is explaining why your chosen format is sufficient for the MVP.

5. Think About Evaluation as a Rubric

Avoid treating a reference solution as the only correct answer. A useful evaluation can look at dimensions such as:

- Requirement understanding
- Class responsibilities
- Coupling / cohesion
- Encapsulation and interfaces
- Appropriate use of abstraction / patterns
- Extensibility when requirements change
- Edge cases and testability
- Quality of explanation

You can score these dimensions, provide structured comments, or combine both. The key is that feedback should point to evidence in the candidate solution.

6. Decide Where AI Helps

A practical design separates deterministic checks from judgment-heavy checks.

Good candidates for deterministic logic	Good candidates for AI
Required fields / structure	Quality of responsibilities
Compilation / tests, if code is supported	Design trade-offs
Known business rules	SOLID / abstraction analysis
Submission state transitions	Candidate explanation analysis
Idempotency / duplicate handling	Improvement suggestions

You do not need a sophisticated AI pipeline. A structured prompt + fixed rubric + stored result can be enough for the prototype.

7. Make AI Feedback More Consistent

One simple approach is to give the model a fixed rubric and require structured output.

For example:

Example evaluation shape

criterion → score → evidence → concern → suggestion → confidence

This is only a starting point. The important idea is to avoid asking the model an unconstrained question such as “Is this a good design?”

8. Keep the LLD of Your Product Meaningful

The core domain may include objects such as Problem, Attempt, Submission, Evaluation, Rubric, and Feedback—but these are examples, not a required model.

Ask of each class:

- What responsibility does it own?
- What behaviour belongs here?
- What does it depend on?
- What is likely to change?
- Do I really need this abstraction?

9. Two Simple Change Tests

Use these to challenge your design:

Change test A

Today the learner submits text. Later the platform supports a class diagram. How much of your domain model changes?

Change test B

Today feedback comes from one evaluator. Later you add a rule-based evaluator or human review. Can you add it without rewriting the practice flow?

You do not need to implement the future features. Your design just needs a sensible place for the variation.

10. Keep Scale Practical

For this assignment, scaling means showing good judgement, not producing a full HLD.

- If AI evaluation is slow, do not block the main submission request unnecessarily.
- Store the submission before evaluation starts so it is not lost on evaluator failure.
- Give evaluation a clear state such as Submitted → Evaluating → Completed / Failed.
- Avoid duplicate processing where a user retries the same request.
- If the product grows, describe the simplest component you would separate first and why.

11. Good Candidate Decisions

We would rather see:

- A simple architecture with clean domain boundaries.
- A clear reason for every important abstraction.
- A useful feedback model instead of a random "AI score".
- A small amount of AI used thoughtfully.
- A prototype that works end-to-end.
- A clear discussion of trade-offs and limitations.

We would be less impressed by:

- Many microservices with no need.
- Lots of UI with weak practice logic.
- Design patterns added only to show pattern knowledge.
- An LLM prompt that simply asks for a 100-point score.
- Large feature lists with little working functionality.

Final reminder

Do not try to solve every possible problem. Solve one clear learner problem well, and use your design to show how the solution can evolve.